

Conclusions

Practical Machine Learning (with R)

UC Berkeley

Spring 2016

3 Questions

- What did you learn that was unexpected and insightful?
- What did you struggle with the most?
- What are you sorry that we didn't get to cover? / What are you going to learn on your own?



Class Evaluation

⇒ Please do it now.



Agenda

⇒ **Administrativa**

- Role Call / 3 Questions
- Class Evaluations

⇒ **Review/Expectations**

- Readings
 - **Forecasting Principals and Practice**
 - Chapters 1 "Getting Started"
 - Chapter 2 "The Forecaster's Toolbox"
 - Chapter 6 "Time Series Decomposition"
 - **APM**
 - Chapter 7.1, 13.2 "Neural Networks"
 - Chapter 7.3, 13.4 "Support Vector Machines"
- Previous Lecture Rewiew

⇒ **New Topics**



Remaining Topics

Today

- Neural Networks
- Deploying Models
- Recommender Patterns
- SVMs





QUESTIONS



NEURAL NETWORKS / ARTIFICIAL NEURAL NETWORKS (ANN)



Analogy to How the Brain Works

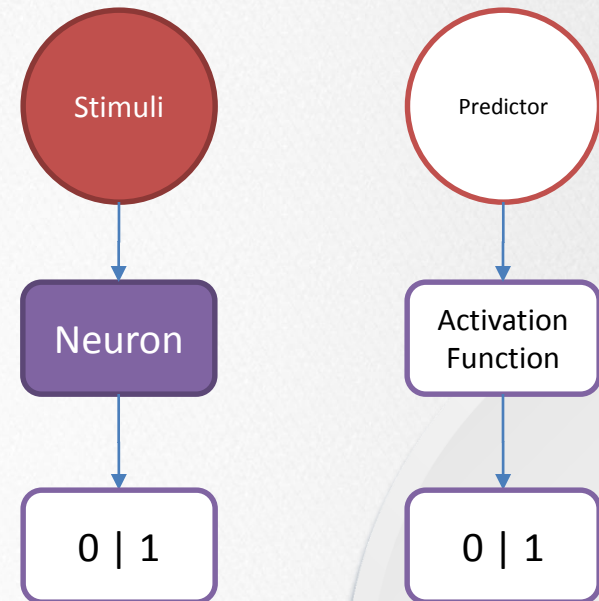
Essential Brain Functions

- Neurons receive signals from stimuli
- Neurons either “fire” or not → 0 | 1
- If fired, signal is propagated.
- Neurons can have multiple inputs and outputs
- Neurons can receive and send signals to other neurons.
 - “Network of neurons” result in cognitive function
- Response is continuous valued



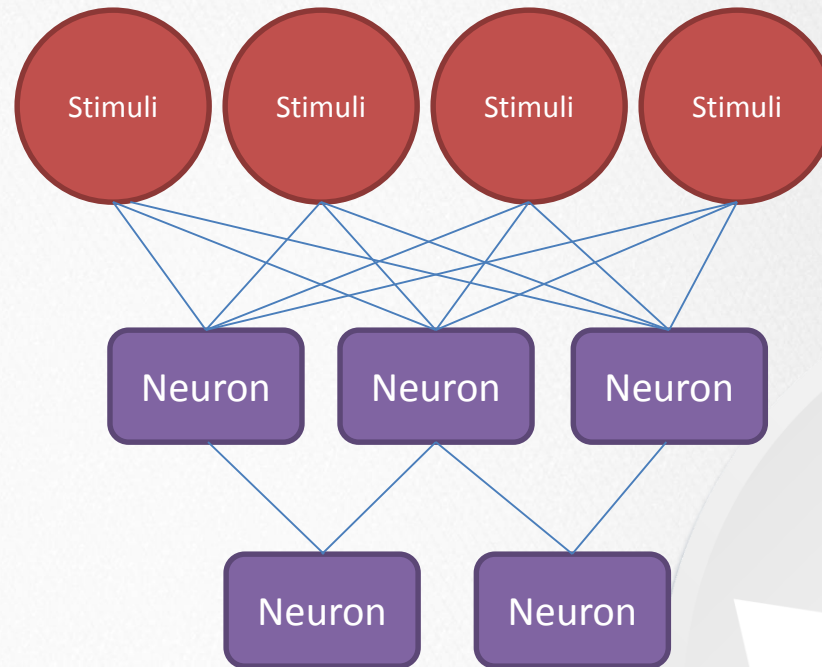
Neuron

- ➔ Neurons receive signals from stimuli
- ➔ Neurons either “fire” or not $\rightarrow 0 \mid 1$



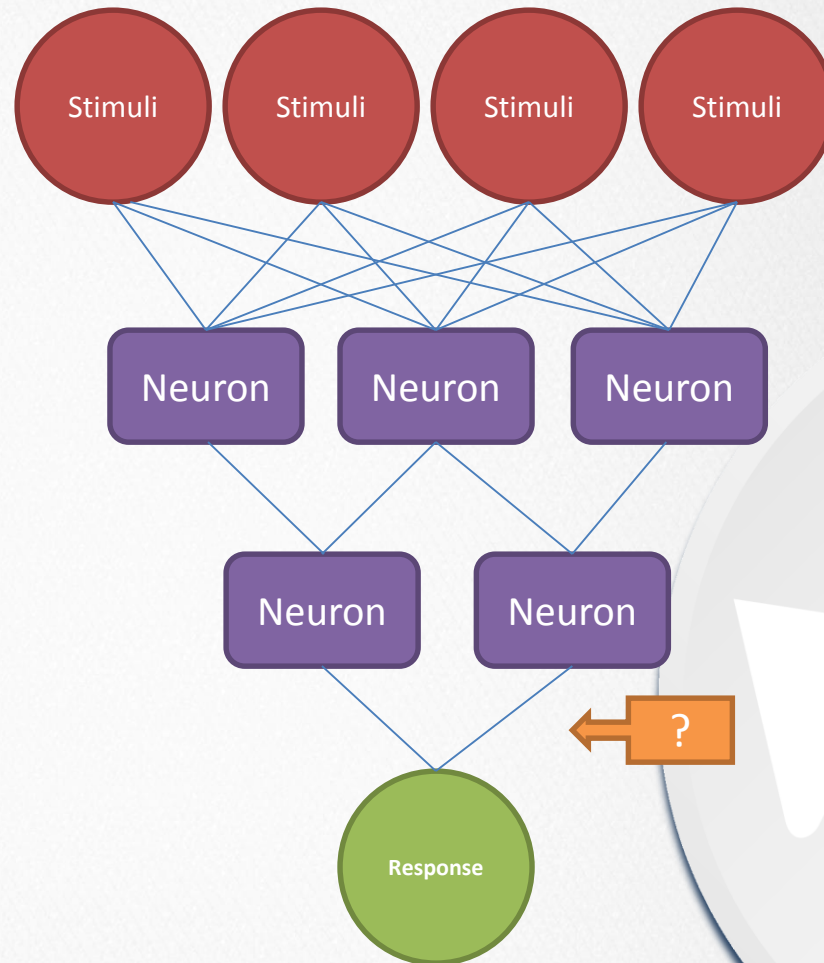
Network

- Neurons can have multiple inputs
- Neurons can send a signal (0,1) to multiple outputs
- Neurons can receive and send signals to other neurons.
 - “Network of neurons” result in cognitive function

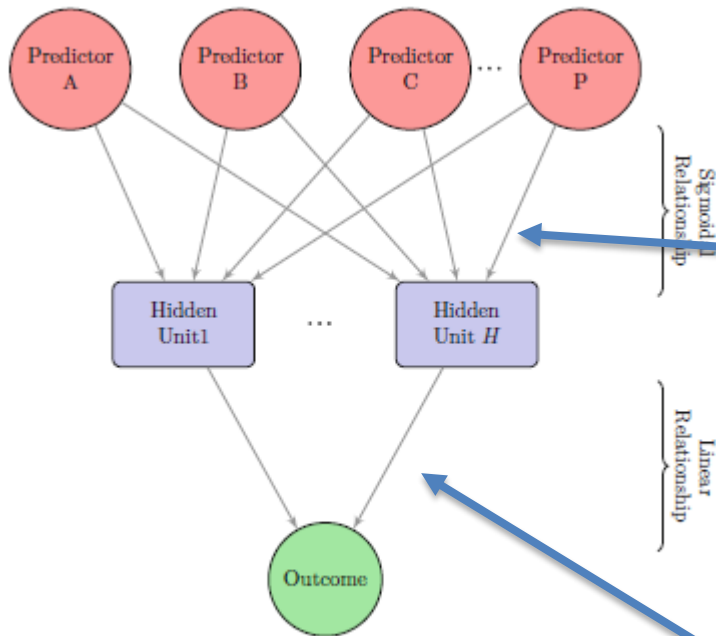


Response

- Response is continuous valued



SINGLE-LAYER FEED-FORWARD NETWORK



Hidden units are comprised of transformation of linear combination of predictors:

$$h(x) = g\left(\beta_0 + \sum_{i=1}^P x_i \beta_i\right)$$

Where $g(x)$ is a *sigmoidal activation* function, resulting in values equal or close to 0 or 1 (often the logistic function).

Outcome is then a regression/classification on the Hidden Units, for example:

$$f(x) = \gamma_0 + \sum_{k=1}^H x_k \gamma_k$$

NB: β and γ are called “weights”.

NEURAL NETWORK AS A LEARNING ALGORITHM

Component	
Loss Function	Various typically: <ul style="list-style-type: none">• RMSE (Regression)• Binary 0-1 Loss (Classification)
Restricted Class of Function	<ul style="list-style-type: none">• Sigmoidal Activation Function (logistic) with• Linear combination of hidden units (linear regression)
Search Methodology	Optimization, typically via Back-propagation

NEURAL NETWORK CLASSIFICATION



NEURAL NETWORKS: VARIATIONS

- Neural Networks can be made more complex by:
 - Adding hidden units
 - Adding hidden layers
 - Allowing back propagation of neuron signals.



NEURAL NETWORK: ADVANTAGES

- ⇒ Performs very well
 - yields highly accurate predictions
 - low signal-to-noise
 - Non-linear response
- ⇒ Easy to add or remove complexity
 - Hidden units
 - Hidden layers
- ⇒ Modelled after brain function



NEURAL NETWORKS: DISADVANTAGES

- Not interpretable → difficult to relate inputs to outcomes
- Requires numeric inputs, typically scaled/centered
- All parameters, β 's and γ 's estimated simultaneously
 - A large number of parameters → requires larger data sets.
$$H * (P + 1) + H + 1$$
 - Training times can be long
 - Numerical solution technique:
 - No guarantee of uniqueness
 - Dependent on initial selections of parameters usually small random
- Tends to overfit

SOLUTIONS TO OVERFITTING

- ➔ **Early stopping:** iterative algorithms for solving for the regression equations can be prematurely halted
- ➔ **Weight Decay:**
add a penalty for large regression coefficients; large values must have a significant effect on the model errors to be tolerated:

$$\sum_{i=1}^n (y_i - f_i(x))^2 .$$

$$0 < \lambda \leq 0.1$$

ASIDE: WEIGHT DECAY

- ➔ This addition of a term to the minimization target is called “shrinkage” or “regularization”
- ➔ Tuning parameter ... Bias Variance Trade-off
- ➔ Can be applied to any technique that directly estimate parameters ... linear regression (ridge L1/lasso L2)



DEEP LEARNING

- ➔ Buzz word: Rebranding of (artificial) neural networks
 - Training speed improvements (Hinton et. al, 2013)
 - Hardware Improvements
- ➔ Requirements
 - Multiple (often) many layers
 - Supervised or unsupervised learning at each layer
 - Each successive layer is a higher level of abstraction

Pixel → Group of pixels → Whisker → Feline Face → House Cat
layer 1 layer 2 layer3

- ➔ Higher abstractions obviate feature engineering ... to a point.
 - Untransformed input features
 - Large success on unsupervised learning, e.g. can we create a model for cats.

Deep Learning

Benefits

- Highly accurate → NN
- Reduced need for manual feature engineering
- Does not necessarily require labelled data
- Exciting area of development ... AI Spring


Limitations

- Computationally expensive
 - long training times or
 - need for specialty hardware
- High number of parameters require lots of data
- Really novel?



Deep Learning

WHAT MAKES DEEP LEARNING DEEP?



Today's Largest Networks

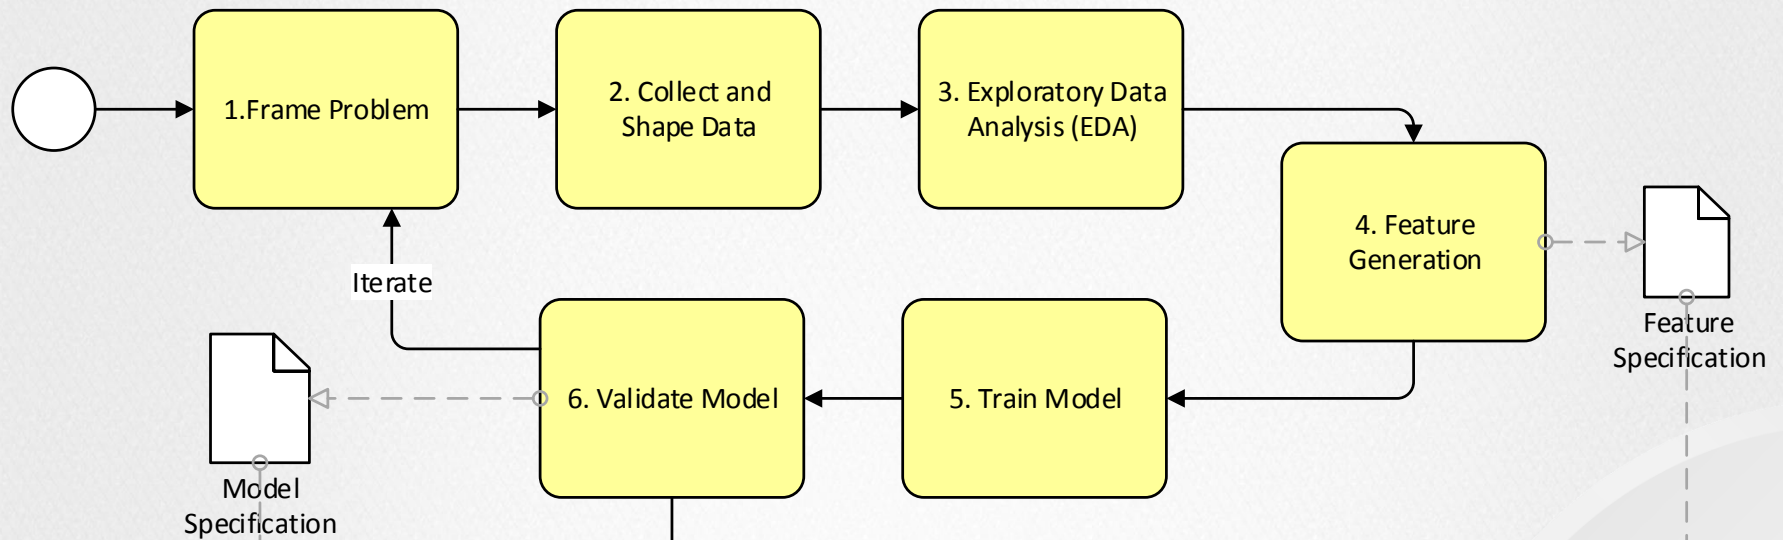
- ~10 layers
- 1B parameters
- 10M images
- ~30 Exaflops
- ~30 GPU days

Human brain has trillions of parameters - only 1,000 more.

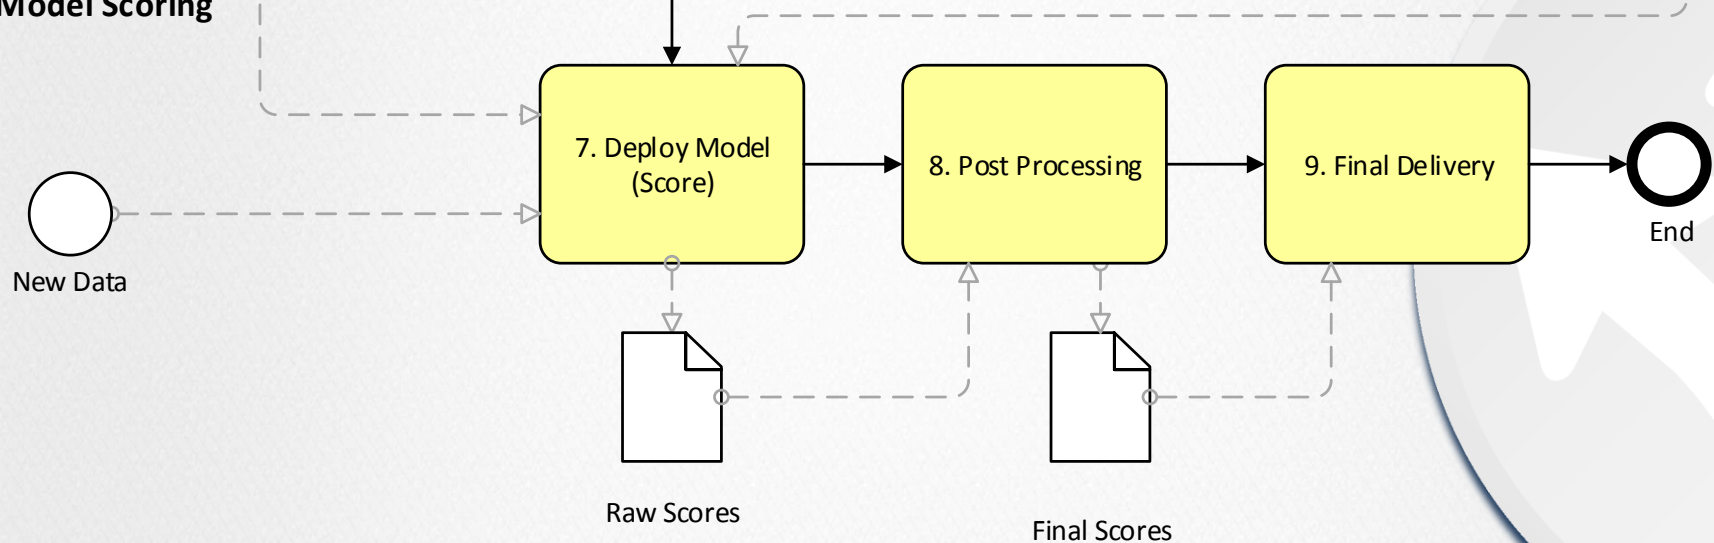
DEPLOYMENT



Model Training



Model Scoring



Source: Decision Patterns

DEPLOYMENT CONSIDERATIONS

- **Access:** How will the model be accessed?
Online? Excel? In Database? In application?
- **Trigger :** What event when scoring occurs? How often?
- **Inputs**
 - Where data to be scored exists, how obtained
 - How much data?
 - What features can be calculated and cached offline?
- **Outputs:** determine the output of the models
 - **Scores :** raw score, class, class probabilities?
 - **Additional data :** e.g. benefit, prescribed actions

- **Logic**

Data → [Feature Specifications] → Features → [Model Specifications] → Final Processing
(online or offline)



MODEL ACCESS AND USAGE PATTERNS

User	Access Point	Deployment Technology	Consideration(s)
Client, Non-technical	Web	Shiny	Interface changes with most model updates
Client, Excel-user	Excel	OpenCPU + VB REST libs	Excel quirks, e.g. Auto calculate
Non-R App. Developer	RESTful API	OpenCPU	Marshalling JSON data to from R
R Developer	Web RESTful API Native R UI R / Rstudio	Shiny OpenCPU Rserve libraries / git / packrat	
Recurrent Process Manual Automated	command line cron	Rscript + optigrab or	Logging, trapping errors, default parameters, notifications (long processes)
Proprietary Application <ul style="list-style-type: none"> Database BI Application 		SQL Server (SSAS), R Oracle Enterprise	Mileage varies

REVIEW: CREATING AN R PACKAGE

A package is just a directory + conventions for locating code, data, documentation and more

To create a package template

```
devtools::create("path/to/mypackage")
```

Packages:

- Save models / feature specs objects to `data/` directory:
 - `save`
- Use `roxygen2` for documentation
 - Uses tags: `@export`, `@param`, `@main`, ...
- Use package metadata for
 - `Author:`, `Date:`, **`Version:`**
- Best practice, use `git tag` to match version number



Shiny – What is it

Interactive web framework for analysis and visualization in R

- Application written entirely in R
- **Reactive** programming
 - Variable have dependency on other values
 - Update values when dependencies change
- Separated concerns
 - **ui/ui.R** (layout, presentation)
 - high-level functions for widgets and layout
 - **server / server.R** (application) : application logic
- Well-documented: shiny.rstudio.com
 - Tutorials, examples, gallery ...
- Integrated with **Rstudio** (`runApp`)
- Uses popular JS libraries for client side reactivity, **DataTable**, **Selectize**, **d3.js**. Search CRAN and Shiny Web sites.



Shiny Resources

- ⇒ shiny.rstudio.com
- ⇒ CRAN
- ⇒ R-bloggers
- ⇒ Stack-overflow



SHINY SCORER EXAMPLE



openCPU

→ Exposes R as a RPC web service:

- `http://myserver/ocpu/path/command/format`

→ Examples

▪ Path:

```
library/stats/  
library/datasets
```

▪ Command

```
R/mean  
data/airquality
```

▪ Format:

```
print | json | csv | tab | md |  
rda | rds | pb | png | pdf | svg
```

→ Example: openCPU example



Shiny vs openCPU

Shiny

- ➔ user-interactive apps
 - single purpose
 - single threaded
- ➔ manages state
 - Remembers uses settings between calls
- ➔ Easy-to do simple things
- ➔ Integration with RStudio

openCPU

- ➔ Web service
 - part of larger application
 - analysis deployed to multiple endpoints
 - Batch
- ➔ Stateless
 - Inherently multithreaded
- ➔ More complex
- ➔ Manual configuration

MISCELLANEOUS TOPICS



Formula Tools

⇒ `i` and `poly`



CUSTOM LOSS FUNCTIONS

- Use in-model training → rarely available
- Use case weights → more weight to are
- Use 2nd optimization → another model



RECOMMENDER SYSTEM



RECOMMENDER SYSTEM

- ⇒ Goal develop a **Netflix/Amazon** style recommendations system to users

Setup

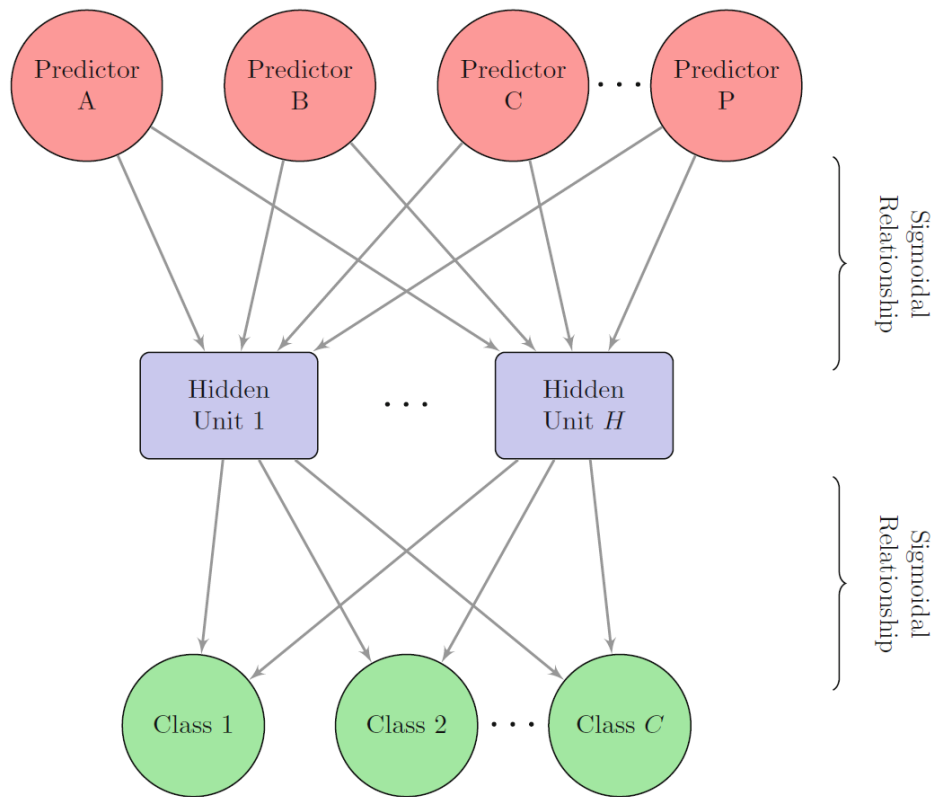
- ⇒ Two disjoint sets: **users** and **products**
- ⇒ Historical interactions:
 - Sales / Rentals
 - Ratings/ Likes /Comments
 - Etc



APPENDIX



NEURAL NETWORKS: CLASSIFICATION




- Add additional responses, one for each class.
- Transform continuous response are transformed by logistic function
- Apply **Softmax** function:
Transform all class probabilities via the soft max function.

$$\hat{p}_\ell^* = \frac{e^{\hat{y}_\ell}}{\sum_{l=1}^C e^{\hat{y}_l}}$$

Maps arbitrary real-values

- so that all values sum to 1
- and are between 0 and 1

Shiny Component



ui.R

- HTML Layout:

- `fluidPage`
- `navBarPage`
- `bootstrapPage`
- `shinydashboard`

- Widgets

- Input: `input*`
- Output: `*Output`

Reactivity:

ui redraws in response to **user inputs** and **server events**

server.R – logic/data/results

- Manage input and output

- `input$name`
- `output$name`

- Manages reactivity

- `observe`, `observeEvent`
- `Isolate`

- Creates data/results:

- `render*`

Reactivity:

server responds to changes in **user inputs**

Shiny: Best Practices

- Best way to learn **shiny** is to do **shiny**
- Shiny applications can get complicated quickly
 - Designed for simple, single-page, single-thread apps
 - Goal: Reduce complexity
- All shiny functions are **camelCase**
 - Separate concerns
 - Keep related things close together
- use `source` as needed
- Separate common elements for both `ui.R` and `server.R` into common file and `source` them or place them in a package



Shiny: Best Practices

→ ui.R

- Program model uses **functions** and **names**
- All ui widgets must have unique name even if identical element is used twice.
- Naming can be confusing → adopt a naming convention.
- For complex applications, separate widgets and layout
 - layout → ui.R
 - widgets → widgets.R

→ Server.R

- Program model uses named lists
 - All inputs accessed by `input` list, must respond appropriately to empty inputs
 - All rendered objects must be placed in `output` list
- Dynamic inputs use `renderUI (server.R)` & `uiOutput (ui.R)`
 - If you use any dynamic inputs, use all dynamic inputs



Shiny: Best Practices

➔ Development

- Develop within Rstudio

```
library(rstudioapi)
```

```
runApp( launch.browser=rstudioapi::viewer ) # or  
options("shiny.launch.browser" = rstudioapi::viewer )
```

- `runApp` can reload with ui.R errors / must be reloaded with server.R errors
- `browser` is your friend!

➔ Open Source free version lacks

- Authentication
- Multiple-threads
- Resource allocation and monitoring

